

Laptop-price-prediction-project

April 4, 2024

```
[159]: import numpy as np
import pandas as pd
df = pd.read_csv("laptop_data.csv")
```

```
[160]: df.head() #To show the first 5 columns of the data set
```

```
[160]: Unnamed: 0 Company  TypeName  Inches  ScreenResolution \
0          0   Apple  Ultrabook   13.3  IPS Panel Retina Display 2560x1600
1          1   Apple  Ultrabook   13.3                1440x900
2          2     HP   Notebook   15.6                Full HD 1920x1080
3          3   Apple  Ultrabook   15.4  IPS Panel Retina Display 2880x1800
4          4   Apple  Ultrabook   13.3  IPS Panel Retina Display 2560x1600
```

```
          Cpu  Ram  Memory \
0  Intel Core i5 2.3GHz  8GB  128GB SSD
1  Intel Core i5 1.8GHz  8GB  128GB Flash Storage
2  Intel Core i5 7200U 2.5GHz  8GB  256GB SSD
3  Intel Core i7 2.7GHz 16GB  512GB SSD
4  Intel Core i5 3.1GHz  8GB  256GB SSD
```

```
          Gpu  OpSys  Weight  Price
0  Intel Iris Plus Graphics 640  macOS  1.37kg  71378.6832
1  Intel HD Graphics 6000  macOS  1.34kg  47895.5232
2  Intel HD Graphics 620  No OS  1.86kg  30636.0000
3  AMD Radeon Pro 455  macOS  1.83kg  135195.3360
4  Intel Iris Plus Graphics 650  macOS  1.37kg  96095.8080
```

```
[161]: df.describe() # This shows the stats of all the numerical values
```

```
[161]: Unnamed: 0  Inches  Price
count  1306.000000  1306.000000  1306.000000
mean    649.690658   15.013017  59912.258499
std     376.885737    1.424616  37304.220229
min       0.000000   10.100000   9270.720000
25%     323.250000   14.000000  31914.720000
50%     649.500000   15.600000  52054.560000
75%     975.750000   15.600000  79373.746800
```

```
max    1302.000000    18.400000  324954.720000
```

```
[162]: df.shape
```

```
[162]: (1306, 12)
```

```
[163]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1306 entries, 0 to 1305
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            1306 non-null   int64
1   Company               1306 non-null   object
2   TypeName              1306 non-null   object
3   Inches               1306 non-null   float64
4   ScreenResolution      1306 non-null   object
5   Cpu                  1306 non-null   object
6   Ram                   1301 non-null   object
7   Memory               1306 non-null   object
8   Gpu                  1306 non-null   object
9   OpSys                1306 non-null   object
10  Weight               1306 non-null   object
11  Price                 1306 non-null   float64
```

```
dtypes: float64(2), int64(1), object(9)
```

```
memory usage: 122.6+ KB
```

DATA CLEANING

```
[164]: df.duplicated().sum() #To find the duplicate values
```

```
[164]: 9
```

```
[165]: df.drop_duplicates(inplace=True) #To drop the duplicate values
```

```
[166]: df.duplicated().sum()
```

```
[166]: 0
```

```
[167]: mode_ram = df['Ram'].mode()[0] # Calculating mode
```

```
[168]: mode_ram
```

```
[168]: '8GB'
```

```
[169]: df['Ram'].fillna(mode_ram, inplace=True)
```

```
[170] : df.isnull().sum() # Null values are replaced by mode
```

```
[170] : Unnamed: 0      0
      Company      0
      TypeName     0
      Inches      0
      ScreenResolution 0
      Cpu         0
      Ram         0
      Memory      0
      Gpu         0
      OpSys       0
      Weight      0
      Price       0
      dtype: int64
```

Data preprocessing

```
[171] : df.drop(columns=['Unnamed: 0'], inplace=True) #Dropping Unnamed: 0 column_
      ↪becuase it is not needed.
```

```
[172] : df.head()
```

```
[172]:
```

	Company	TypeName	Inches	ScreenResolution	
0	Apple	Ultrabook	13.3	IPS Panel Retina Display	2560x1600
1	Apple	Ultrabook	13.3		1440x900
2	HP	Notebook	15.6	Full HD	1920x1080
3	Apple	Ultrabook	15.4	IPS Panel Retina Display	2880x1800
4	Apple	Ultrabook	13.3	IPS Panel Retina Display	2560x1600

		Cpu	Ram	Memory	
0	Intel Core i5	2.3GHz	8GB	128GB SSD	
1	Intel Core i5	1.8GHz	8GB	128GB Flash Storage	
2	Intel Core i5 7200U	2.5GHz	8GB	256GB SSD	
3	Intel Core i7	2.7GHz	16GB	512GB SSD	
4	Intel Core i5	3.1GHz	8GB	256GB SSD	

		Gpu	OpSys	Weight	Price
0	Intel Iris Plus Graphics	640	macOS	1.37kg	71378.6832
1	Intel HD Graphics	6000	macOS	1.34kg	47895.5232
2	Intel HD Graphics	620	No OS	1.86kg	30636.0000
3	AMD Radeon Pro	455	macOS	1.83kg	135195.3360
4	Intel Iris Plus Graphics	650	macOS	1.37kg	96095.8080

```
[173] : df['Ram'] = df['Ram'].str.replace('GB',"") #To remove 'GB' from Ram column
```

```
[174] : df['Weight'] = df['Weight'].str.replace('kg',"") #To remove 'kg' from Ram column
```

```
[175] : df.head()
```

```
[175]:
```

	Company	TypeName	Inches				ScreenResolution	\
0	Apple	Ultrabook	13.3	IPS	Panel	Retina	Display	2560x1600
1	Apple	Ultrabook	13.3					1440x900
2	HP	Notebook	15.6				Full HD	1920x1080
3	Apple	Ultrabook	15.4	IPS	Panel	Retina	Display	2880x1800
4	Apple	Ultrabook	13.3	IPS	Panel	Retina	Display	2560x1600

			Cpu	Ram			Memory	\
0		Intel Core i5	2.3GHz	8			128GB SSD	
1		Intel Core i5	1.8GHz	8	128GB	Flash Storage		
2	Intel	Core i5 7200U	2.5GHz	8			256GB SSD	
3		Intel Core i7	2.7GHz	16			512GB SSD	
4		Intel Core i5	3.1GHz	8			256GB SSD	

			Gpu	OpSys	Weight		Price
0	Intel	Iris Plus Graphics	640	macOS	1.37		71378.6832
1		Intel HD Graphics	6000	macOS	1.34		47895.5232
2		Intel HD Graphics	620	No OS	1.86		30636.0000
3		AMD Radeon Pro	455	macOS	1.83		135195.3360
4	Intel	Iris Plus Graphics	650	macOS	1.37		96095.8080

```
[176] : df['Ram'] = df['Ram'].astype('int32') #Changing the datatypes
```

```
[177] : df['Weight'] = df['Weight'].astype('float32')
```

```
[178] : df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1297 entries, 0 to 1305
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Company                1297 non-null   object
1   TypeName               1297 non-null   object
2   Inches                 1297 non-null   float64
3   ScreenResolution       1297 non-null   object
4   Cpu                    1297 non-null   object
5   Ram                    1297 non-null   int32
6   Memory                 1297 non-null   object
7   Gpu                    1297 non-null   object
8   OpSys                  1297 non-null   object
9   Weight                 1297 non-null   float32
10  Price                  1297 non-null   float64
dtypes: float32(1), float64(2), int32(1), object(7)
memory usage: 111.5+ KB
```

```
[179] : df.describe() # Found outliers in Weight
```

```
[179]:
```

	Inches	Ram	Weight	Price
count	1297.000000	1297.000000	1297.000000	1297.000000
mean	15.013647	8.390131	2.806654	59924.888938
std	1.425709	5.091038	27.718493	37296.015690
min	10.100000	2.000000	0.690000	9270.720000
25%	14.000000	4.000000	1.500000	31914.720000
50%	15.600000	8.000000	2.040000	52054.560000
75%	15.600000	8.000000	2.300000	79387.200000
max	18.400000	64.000000	1000.000000	324954.720000

Removing outliers

```
[180] : Q1 = df.Weight.quantile(0.0) # Choosing 0 after trial and testing of range
      : Q3 = df.Weight.quantile(0.75)
```

```
[181] : Q1
```

```
[181]: 0.6899999976158142
```

```
[182] : Q3
```

```
[182]: 2.299999952316284
```

```
[183] : IQR = Q3-Q1
```

```
[184] : IQR
```

```
[184]: 1.60999995470047
```

```
[185] : lower_limit = Q1 - 1.5*IQR
      : upper_limit = Q3 + 1.5*IQR
      : lower_limit, upper_limit
```

```
[185]: (-1.7249999344348907, 4.714999884366989)
```

```
[186] : df[(df.Weight<lower_limit) | (df.Weight>upper_limit)] #To find the outliers
```

```
[186] :
```

	Company	Type	Name	Inches	ScreenResolution	Cpu \
306	Lenovo	Notebook		15.6	1366x768	Intel Core i5 7200U 2.5GHz

	Ram	Memory	Gpu	OpSys	Weight	Price
306	8	2TB HDD	Nvidia GeForce 940MX	No OS	1000.0	29250.72

```
[187] : df = df[(df.Weight>lower_limit)&(df.Weight<upper_limit)] # To remove the
      : outliers
```

```
df.describe()
```

```
[187]:
```

	Inches	Ram	Weight	Price
count	1296.000000	1296.000000	1296.000000	1296.000000
mean	15.013194	8.390432	2.037215	59948.557278
std	1.426166	5.092992	0.665834	37300.667342
min	10.100000	2.000000	0.690000	9270.720000
25%	14.000000	4.000000	1.500000	31914.720000
50%	15.600000	8.000000	2.040000	52107.840000
75%	15.600000	8.000000	2.300000	79424.496000
max	18.400000	64.000000	4.700000	324954.720000

Exploratory Data Analysis

```
[188] : import seaborn as sns
import matplotlib.pyplot as plt
```

```
[189] : sns.distplot(df['Price'])
```

<ipython-input-189-87e11caeb2c4>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df['Price'])
```

```
[189] : <Axes: xlabel='Price', ylabel='Density'>
```




```
[190] : df['Company'].value_counts().plot(kind='bar') #To show the count of laptop of  
         each brand.
```

```
[190] : <Axes: xlabel='Company'>
```




```
[191] : sns.barplot(x=df['Company'], y=df['Price']) # To find which laptops are costly,  
        and which one are budget laptops.  
        plt.xticks(rotation='vertical')  
        plt.show()
```





```
[192] : df['TypeName'].value_counts().plot(kind='bar') #To find the count of
```

```
[192] : <Axes: xlabel='TypeName'>
```



```
[193] : sns.barplot(x=df['TypeName'], y=df['Price']) # To find which type of laptops  
        are costly and which one are budget laptops.  
        plt.xticks(rotation='vertical')  
        plt.show()
```



```
[194]: sns.distplot(df['Inches'])
```

<ipython-input-194-51888cb550e6>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df['Inches'])
```

```
[194]: <Axes: xlabel='Inches', ylabel='Density'>
```



```
[195] : sns.scatterplot(x=df['Inches'], y=df['Price']) #To check the relation between  
↪Inches and price
```

```
[195] : <Axes: xlabel='Inches', ylabel='Price'>
```



```
[196] : df['Touchscreen']=df['ScreenResolution'].apply(lambda x:1 if 'Touchscreen' in x_
        ↪else 0) #To create a new column named 'Touchscreen' if the word is present.
```

```
[197] : df.head()
```

```
[197]:
```

	Company	TypeName	Inches	ScreenResolution \
0	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600
1	Apple	Ultrabook	13.3	1440x900
2	HP	Notebook	15.6	Full HD 1920x1080
3	Apple	Ultrabook	15.4	IPS Panel Retina Display 2880x1800
4	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600


	Cpu	Ram	Memory \
0	Intel Core i5 2.3GHz	8	128GB SSD
1	Intel Core i5 1.8GHz	8	128GB Flash Storage
2	Intel Core i5 7200U 2.5GHz	8	256GB SSD
3	Intel Core i7 2.7GHz	16	512GB SSD
4	Intel Core i5 3.1GHz	8	256GB SSD

	Gpu	OpSys	Weight	Price	Touchscreen
0	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	0

1	Intel HD Graphics 6000	macOS	1.34	47895.5232	0
2	Intel HD Graphics 620	No OS	1.86	30636.0000	0
3	AMD Radeon Pro 455	macOS	1.83	135195.3360	0
4	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	0


```
[198]: df['Touchscreen'].value_counts().plot(kind='bar')
```

```
[198]: <Axes: xlabel='Touchscreen'>
```



```
[199]: sns.barplot(x=df['Touchscreen'], y=df['Price']) #To show the relation between
↪ Touchscreen and Price
```


```
[199]: <Axes: xlabel='Touchscreen', ylabel='Price'>
```



```
[200] : df['IPS'] = df['ScreenResolution'].apply(lambda x:1 if 'IPS' in x else 0)
```

```
[201] : sns.barplot(x=df['IPS'], y=df['Price']) #To show the relation between IPS and Price
```

```
[201] : <Axes: xlabel='IPS', ylabel='Price'>
```



```
[202] : temp = df['ScreenResolution'].str.split('x',n=1,expand=True)
```

'x': This is the delimiter used for splitting the strings.

n=1: This parameter indicates that only the first occurrence of the delimiter should be used for splitting. This means if there are multiple occurrences of 'x' in a string, only the first one will be considered for splitting.

expand=True: This parameter tells pandas to expand the split strings into separate columns. Each part of the split string will be placed into a separate column in the resulting DataFrame.

```
[203] : df['X_res'] = temp[0]
        df['Y_res'] = temp[1]
```

```
[204] : df.head(10)
```

```
[204]: Company  TypeName  Inches  ScreenResolution \
0   Apple  Ultrabook   13.3  IPS Panel Retina Display 2560x1600
1   Apple  Ultrabook   13.3                1440x900
2    HP   Notebook   15.6                Full HD 1920x1080
3   Apple  Ultrabook   15.4  IPS Panel Retina Display 2880x1800
4   Apple  Ultrabook   13.3  IPS Panel Retina Display 2560x1600
```


5	Acer	Notebook	15.6				1366x768
6	Apple	Ultrabook	15.4	IPS Panel	Retina Display	2880x1800	
7	Apple	Ultrabook	13.3			1440x900	
8	Asus	Ultrabook	14.0			Full HD 1920x1080	
9	Acer	Ultrabook	14.0	IPS Panel	Full HD	1920x1080	

		Cpu	Ram	Memory	\
0	Intel Core i5	2.3GHz	8	128GB SSD	
1	Intel Core i5	1.8GHz	8	128GB Flash Storage	
2	Intel Core i5 7200U	2.5GHz	8	256GB SSD	
3	Intel Core i7	2.7GHz	16	512GB SSD	
4	Intel Core i5	3.1GHz	8	256GB SSD	
5	AMD A9-Series 9420	3GHz	4	500GB HDD	
6	Intel Core i7	2.2GHz	16	256GB Flash Storage	
7	Intel Core i5	1.8GHz	8	256GB Flash Storage	
8	Intel Core i7 8550U	1.8GHz	16	512GB SSD	
9	Intel Core i5 8250U	1.6GHz	8	256GB SSD	

		Gpu	OpSys	Weight	Price	Touchscreen	\
0	Intel Iris Plus Graphics	640	macOS	1.37	71378.6832	0	
1	Intel HD Graphics	6000	macOS	1.34	47895.5232	0	
2	Intel HD Graphics	620	No OS	1.86	30636.0000	0	
3	AMD Radeon Pro	455	macOS	1.83	135195.3360	0	
4	Intel Iris Plus Graphics	650	macOS	1.37	96095.8080	0	
5	AMD Radeon R5		Windows 10	2.10	21312.0000	0	
6	Intel Iris Pro Graphics		Mac OS X	2.04	114017.6016	0	
7	Intel HD Graphics	6000	macOS	1.34	61735.5360	0	
8	Nvidia GeForce MX150		Windows 10	1.30	79653.6000	0	
9	Intel UHD Graphics	620	Windows 10	1.60	41025.6000	0	

	IPS		X_res	Y_res
0	1	IPS Panel Retina Display	2560	1600
1	0		1440	900
2	0	Full HD	1920	1080
3	1	IPS Panel Retina Display	2880	1800
4	1	IPS Panel Retina Display	2560	1600
5	0		1366	768
6	1	IPS Panel Retina Display	2880	1800
7	0		1440	900
8	0	Full HD	1920	1080
9	1	IPS Panel Full HD	1920	1080

```
[205]: df['X_res'].str.replace(',','').str.findall(r'(\d+\.?\d+)') # The output is an array
```

```
[205]: 0    [2560]
       1    [1440]
```

```

2      [1920]
3      [2880]
4      [2560]
...
1301   [1920]
1302   [3200]
1303   [1366]
1304   [1366]
1305   [1366]

```

Name: X_res, Length: 1296, dtype: object

Putting it all together, the pattern `(+.\?+)` matches sequences of one or more digits, optionally followed by a dot, and then followed by one or more digits. This pattern is commonly used to extract numerical values (possibly with decimal points) from strings.

```

[206] : df['X_res']=df['X_res'].str.replace(',','').str.findall(r'(\d+\.? \d+)').
        ↪ apply(lambda x:x[0]) #Applied Lambda function to return the 0th value of the_
        ↪ array

```

```

[207] : df.head(10)

```

```

[207]: Company  TypeName  Inches  ScreenResolution \
0   Apple  Ultrabook   13.3  IPS Panel Retina Display 2560x1600
1   Apple  Ultrabook   13.3                1440x900
2    HP   Notebook   15.6                Full HD 1920x1080
3   Apple  Ultrabook   15.4  IPS Panel Retina Display 2880x1800
4   Apple  Ultrabook   13.3  IPS Panel Retina Display 2560x1600
5   Acer  Notebook   15.6                1366x768
6   Apple  Ultrabook   15.4  IPS Panel Retina Display 2880x1800
7   Apple  Ultrabook   13.3                1440x900
8   Asus  Ultrabook   14.0                Full HD 1920x1080
9   Acer  Ultrabook   14.0                IPS Panel Full HD 1920x1080

```

```

          Cpu  Ram  Memory \
0      Intel Core i5 2.3GHz    8    128GB SSD
1      Intel Core i5 1.8GHz    8  128GB Flash Storage
2  Intel Core i5 7200U 2.5GHz    8    256GB SSD
3      Intel Core i7 2.7GHz   16    512GB SSD
4      Intel Core i5 3.1GHz    8    256GB SSD
5    AMD A9-Series 9420 3GHz    4    500GB HDD
6      Intel Core i7 2.2GHz   16  256GB Flash Storage
7      Intel Core i5 1.8GHz    8  256GB Flash Storage
8  Intel Core i7 8550U 1.8GHz   16    512GB SSD
9  Intel Core i5 8250U 1.6GHz    8    256GB SSD

```

```

          Gpu  OpSys  Weight  Price  Touchscreen \
0  Intel Iris Plus Graphics 640  macOS  1.37  71378.6832  0

```

1	Intel HD Graphics 6000	macOS	1.34	47895.5232	0
2	Intel HD Graphics 620	No OS	1.86	30636.0000	0
3	AMD Radeon Pro 455	macOS	1.83	135195.3360	0
4	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	0
5	AMD Radeon R5	Windows 10	2.10	21312.0000	0
6	Intel Iris Pro Graphics	Mac OS X	2.04	114017.6016	0
7	Intel HD Graphics 6000	macOS	1.34	61735.5360	0
8	Nvidia GeForce MX150	Windows 10	1.30	79653.6000	0
9	Intel UHD Graphics 620	Windows 10	1.60	41025.6000	0

	IPS	X_res	Y_res
0	1	2560	1600
1	0	1440	900
2	0	1920	1080
3	1	2880	1800
4	1	2560	1600
5	0	1366	768
6	1	2880	1800
7	0	1440	900
8	0	1920	1080
9	1	1920	1080

```
[208] : df['X_res'] = df['X_res'].astype('int32')
df['Y_res'] = df['Y_res'].astype('int32')
```

```
[209] : df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1296 entries, 0 to 1305
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Company                1296 non-null   object
1   TypeName               1296 non-null   object
2   Inches                 1296 non-null   float64
3   ScreenResolution       1296 non-null   object
4   Cpu                    1296 non-null   object
5   Ram                    1296 non-null   int32
6   Memory                 1296 non-null   object
7   Gpu                    1296 non-null   object
8   OpSys                  1296 non-null   object
9   Weight                 1296 non-null   float32
10  Price                  1296 non-null   float64
11  Touchscreen            1296 non-null   int64
12  IPS                    1296 non-null   int64
13  X_res                  1296 non-null   int32
14  Y_res                  1296 non-null   int32
```

```
dtypes: float32(1), float64(2), int32(3), int64(2), object(7)
memory usage: 141.8+ KB
```

```
[210] : df.shape
```

```
[210]: (1296, 15)
```

PPI is a measure of pixel density or resolution of a computer screen, television screen or other display device.

```
[211] : df['ppi'] = (((df['X_res']**2)+(df['Y_res']**2)**0.5/df['Inches'])).
         ↳astype('float')
```

```
[212] : df.drop(columns=['ScreenResolution','X_res','Y_res', 'Inches'], inplace=True)
```

```
[213] : df.head()
```

```
[213]:
```

	Company	TypeName		Cpu	Ram	Memory \
0	Apple	Ultrabook	Intel Core i5	2.3GHz	8	128GB SSD
1	Apple	Ultrabook	Intel Core i5	1.8GHz	8	128GB Flash Storage
2	HP	Notebook	Intel Core i5 7200U	2.5GHz	8	256GB SSD
3	Apple	Ultrabook	Intel Core i7	2.7GHz	16	512GB SSD
4	Apple	Ultrabook	Intel Core i5	3.1GHz	8	256GB SSD

			Gpu	OpSys	Weight	Price	Touchscreen	IPS \
0	Intel	Iris Plus Graphics	640	macOS	1.37	71378.6832	0	1
1		Intel HD Graphics	6000	macOS	1.34	47895.5232	0	0
2		Intel HD Graphics	620	No OS	1.86	30636.0000	0	0
3		AMD Radeon Pro	455	macOS	1.83	135195.3360	0	1
4	Intel	Iris Plus Graphics	650	macOS	1.37	96095.8080	0	1


```
      ppi
0 226.983005
1 127.677940
2 141.211998
3 220.534624
4 226.983005
```

```
[214] : df['Cpu']
```

```
[214] : 0          Intel Core i5 2.3GHz
      1          Intel Core i5 1.8GHz
      2          Intel Core i5 7200U 2.5GHz
```

```

3          Intel Core i7 2.7GHz
4          Intel Core i5 3.1GHz
...
1301       Intel Core i7 6500U 2.5GHz
1302       Intel Core i7 6500U 2.5GHz
1303 Intel Celeron Dual Core N3050 1.6GHz
1304       Intel Core i7 6500U 2.5GHz
1305 Intel Celeron Dual Core N3050 1.6GHz
Name: Cpu, Length: 1296, dtype: object

```

```
[215] : df['Cpu Name']=df['Cpu'].apply(lambda x:" ".join(x.split()[0:3])) # To get
↳the first three words
```

```
[216] : df['Cpu Name']
```

```

[216] : 0          Intel Core i5
        1          Intel Core i5
        2          Intel Core i5
        3          Intel Core i7
        4          Intel Core i5
...
1301       Intel Core i7
1302       Intel Core i7
1303 Intel Celeron Dual
1304       Intel Core i7
1305 Intel Celeron Dual
Name: Cpu Name, Length: 1296, dtype: object

```

```

[217] : def fetch_processor(text):
        if text == 'Intel Core i7' or text == 'Intel Core i3' or text == 'Intel Core_
↳i5':
            return text
        else:
            if text.split()[0] == 'Intel':
                return 'Other Intel Processor'
            else:
                return 'AMD Processor'

```

```
[218] : df['Cpu brand'] = df['Cpu Name'].apply(fetch_processor)
```

```
[219] : df.head()
```

```

[219]: Company  TypeName          Cpu  Ram      Memory \
0   Apple  Ultrabook   Intel Core i5 2.3GHz    8      128GB SSD
1   Apple  Ultrabook   Intel Core i5 1.8GHz    8  128GB Flash Storage
2    HP    Notebook Intel Core i5 7200U 2.5GHz    8      256GB SSD
3   Apple  Ultrabook   Intel Core i7 2.7GHz   16      512GB SSD

```


4	Apple	Ultrabook	Intel	Core i5	3.1GHz	8	256GB SSD
---	-------	-----------	-------	---------	--------	---	-----------

				Gpu	OpSys	Weight	Price	Touchscreen	IPS	\
0	Intel	Iris Plus Graphics	640	macOS	1.37	71378.6832	0	1		
1		Intel HD Graphics	6000	macOS	1.34	47895.5232	0	0		
2		Intel HD Graphics	620	No OS	1.86	30636.0000	0	0		
3		AMD Radeon Pro	455	macOS	1.83	135195.3360	0	1		
4	Intel	Iris Plus Graphics	650	macOS	1.37	96095.8080	0	1		


	ppi	Cpu Name	Cpu brand
0	226.983005	Intel Core i5	Intel Core i5
1	127.677940	Intel Core i5	Intel Core i5
2	141.211998	Intel Core i5	Intel Core i5
3	220.534624	Intel Core i7	Intel Core i7
4	226.983005	Intel Core i5	Intel Core i5

```
[220] : df['Cpu brand'].value_counts().plot(kind='bar')
```

```
[220] : <Axes: xlabel='Cpu brand'>
```



```
[221] : sns.barplot(x=df['Cpu brand'], y=df['Price'])  
plt.xticks(rotation='vertical')  
plt.show()
```



```
[222] : df.drop(columns=['Cpu','Cpu Name'], inplace=True) # Dropping the CPU and Cpu_
        Name column because we have created a new column Cpu brand
```

```
[223] : df.drop(columns=['Memory'], inplace=True)
```

```
[224] : df.head()
```

```
[224]: Company  TypeName  Ram      Gpu  OpSys  Weight \
0   Apple  Ultrabook    8  Intel Iris Plus Graphics 640  macOS    1.37
1   Apple  Ultrabook    8      Intel HD Graphics 6000  macOS    1.34
2    HP    Notebook    8      Intel HD Graphics 620  No OS    1.86
3   Apple  Ultrabook   16      AMD Radeon Pro 455  macOS    1.83
4   Apple  Ultrabook    8  Intel Iris Plus Graphics 650  macOS    1.37
```


	Price	Touchscreen	IPS	ppi	Cpu brand
0	71378.6832	0	1	226.983005	Intel Core i5
1	47895.5232	0	0	127.677940	Intel Core i5
2	30636.0000	0	0	141.211998	Intel Core i5
3	135195.3360	0	1	220.534624	Intel Core i7
4	96095.8080	0	1	226.983005	Intel Core i5

```
[225] : df['Gpu'].unique()
```

```
[225] : array(['Intel Iris Plus Graphics 640', 'Intel HD Graphics 6000',
              'Intel HD Graphics 620', 'AMD Radeon Pro 455',
              'Intel Iris Plus Graphics 650', 'AMD Radeon R5',
              'Intel Iris Pro Graphics', 'Nvidia GeForce MX150',
              'Intel UHD Graphics 620', 'Intel HD Graphics 520',
              'AMD Radeon Pro 555', 'AMD Radeon R5 M430',
              'Intel HD Graphics 615', 'AMD Radeon Pro 560',
              'Nvidia GeForce 940MX', 'Intel HD Graphics 400',
              'Nvidia GeForce GTX 1050', 'AMD Radeon R2', 'AMD Radeon 530',
              'Nvidia GeForce 930MX', 'Intel HD Graphics',
              'Intel HD Graphics 500', 'Nvidia GeForce 930MX ',
              'Nvidia GeForce GTX 1060', 'Nvidia GeForce 920MX',
              'AMD Radeon R4 Graphics', 'AMD Radeon 520',
              'Nvidia GeForce GTX 1070', 'Nvidia GeForce GTX 1050 Ti',
              'Nvidia GeForce MX130', 'AMD R4 Graphics',
              'Nvidia GeForce GTX 940MX', 'AMD Radeon RX 560',
              'Nvidia GeForce 920M', 'AMD Radeon R7 M445', 'AMD Radeon RX 550',
              'Nvidia GeForce GTX 1050M', 'Intel HD Graphics 515',
              'AMD Radeon R5 M420', 'Intel HD Graphics 505',
              'Nvidia GTX 980 SLI', 'AMD R17M-M1-70', 'Nvidia GeForce 150MX',
              'Nvidia GeForce GTX 1080', 'Nvidia Quadro M1200',
              'Nvidia GeForce 920MX ', 'Nvidia GeForce GTX 950M',
              'AMD FirePro W4190M ', 'Nvidia GeForce GTX 980M',
              'Intel Iris Graphics 550', 'AMD Radeon RX 580',
              'Nvidia GeForce 930M', 'Intel HD Graphics 630',
              'AMD Radeon R5 430', 'Nvidia GeForce GTX 940M',
              'Intel HD Graphics 510', 'Intel HD Graphics 405',
              'AMD Radeon RX 540', 'Nvidia GeForce GT 940MX',
              'AMD FirePro W5130M', 'Nvidia Quadro M2200M', 'AMD Radeon R4',
              'Nvidia Quadro M620', 'AMD Radeon R7 M460',
              'Intel HD Graphics 530', 'Nvidia GeForce GTX 965M',
              'Nvidia GeForce GTX1080', 'Nvidia GeForce GTX1050 Ti',
              'Nvidia GeForce GTX 960M', 'AMD Radeon R2 Graphics',
              'Nvidia Quadro M620M', 'Nvidia GeForce GTX 970M',
              'Nvidia GeForce GTX 960<U+039C>', 'Intel Graphics 620',
              'Nvidia GeForce GTX 960', 'AMD Radeon R5 520',
              'AMD Radeon R7 M440', 'AMD Radeon R7', 'Nvidia Quadro M520M',
              'Nvidia Quadro M2200', 'Nvidia Quadro M2000M',
```

```
'Intel HD Graphics 540', 'Nvidia Quadro M1000M', 'AMD Radeon 540',
'Nvidia GeForce GTX 1070M', 'Nvidia GeForce GTX1060',
'Intel HD Graphics 5300', 'AMD Radeon R5 M420X',
'AMD Radeon R7 Graphics', 'Nvidia GeForce 920',
'Nvidia GeForce 940M', 'Nvidia GeForce GTX 930MX',
'AMD Radeon R7 M465', 'AMD Radeon R3', 'Nvidia GeForce GTX 1050Ti',
'AMD Radeon R7 M365X', 'Intel Iris Graphics 540',
'AMD Radeon R9 M385', 'Intel HD Graphics 620 ',
'Nvidia Quadro 3000M', 'Nvidia GeForce GTX 980 ',
'AMD Radeon R5 M330', 'AMD FirePro W4190M', 'AMD FirePro W6150M',
'AMD Radeon R5 M315', 'Nvidia Quadro M500M', 'AMD Radeon R7 M360',
'Nvidia Quadro M3000M', 'Nvidia GeForce 960M', 'ARM Mali T860 MP4'],
dtype=object)
```

```
[226] : df['Gpu'].apply(lambda x:x.split()[0])
```

```
[226]: 0      Intel
      1      Intel
      2      Intel
      3      AMD
      4      Intel
      ...
     1301     Intel
     1302     Intel
     1303     Intel
     1304      AMD
     1305     Intel
      Name: Gpu, Length: 1296, dtype: object
```

```
[227] : df['Gpu brand'] = df['Gpu'].apply(lambda x:x.split()[0])
```

```
[228] : df.head()
```

```
[228]: Company  TypeName  Ram      Gpu  OpSys  Weight \
0  Apple  Ultrabook    8 Intel Iris Plus Graphics 640  macOS    1.37
1  Apple  Ultrabook    8      Intel HD Graphics 6000  macOS    1.34
2    HP  Notebook    8      Intel HD Graphics 620  No OS    1.86
3  Apple  Ultrabook   16      AMD Radeon Pro 455  macOS    1.83
4  Apple  Ultrabook    8 Intel Iris Plus Graphics 650  macOS    1.37

      Price  Touchscreen  IPS      ppi      Cpu brand  Gpu brand
0  71378.6832                226.983005 Intel Core i5      Intel
1  47895.5232                127.677940 Intel Core i5      Intel
2  30636.0000                141.211998 Intel Core i5      Intel
3  135195.3360               220.534624 Intel Core i7      AMD
4   96095.8080               226.983005 Intel Core i5      Intel
```

```
[229] : df.drop(columns=['Gpu'], inplace=True) # Drop the GPU column
```

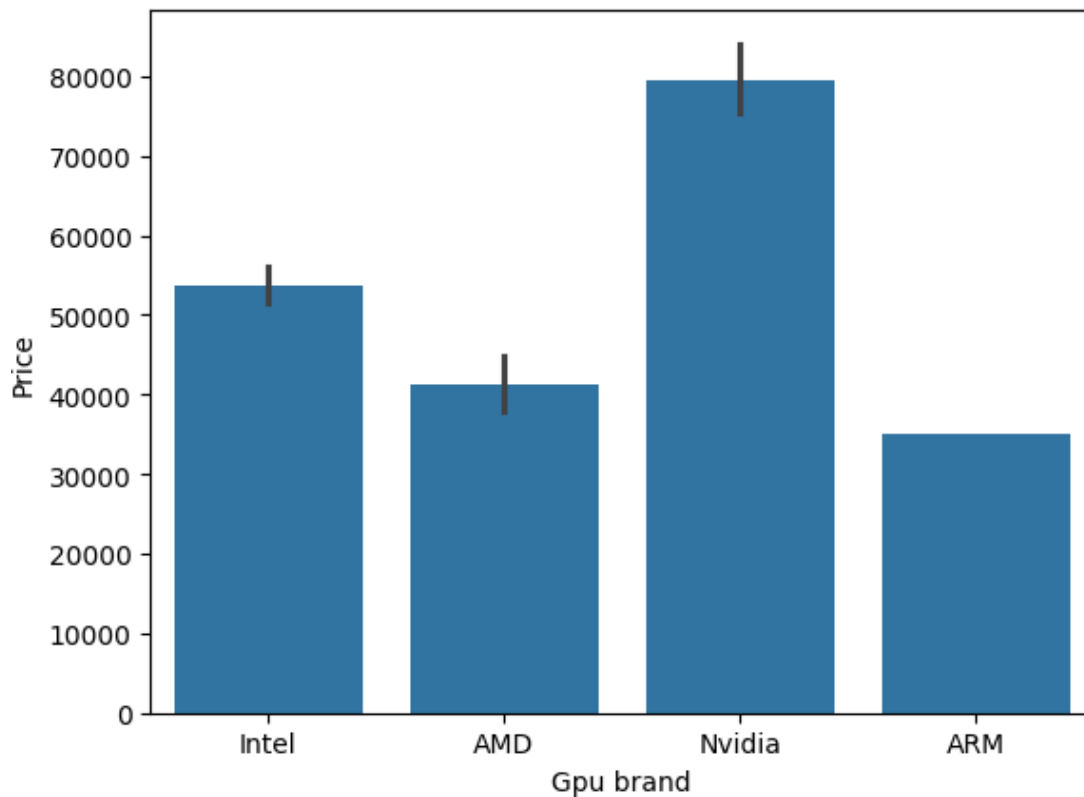
```
[230] : df.head()
```

```
[230] : Company TypeName Ram OpSys Weight Price Touchscreen IPS \
0 Apple Ultrabook 8 macOS 1.37 71378.6832 0 1
1 Apple Ultrabook 8 macOS 1.34 47895.5232 0 0
2 HP Notebook 8 No OS 1.86 30636.0000 0 0
3 Apple Ultrabook 16 macOS 1.83 135195.3360 0 1
4 Apple Ultrabook 8 macOS 1.37 96095.8080 0 1
```

```
ppi Cpu brand Gpu brand
0 226.983005 Intel Core i5 Intel
1 127.677940 Intel Core i5 Intel
2 141.211998 Intel Core i5 Intel
3 220.534624 Intel Core i7 AMD
4 226.983005 Intel Core i5 Intel
```

```
[231] : sns.barplot(x=df['Gpu brand'], y=df['Price'])
```

```
[231] : <Axes: xlabel='Gpu brand', ylabel='Price'>
```



```
[232] : df['OpSys'].value_counts()
```

```
[232] : OpSys
Windows 10      1068
No OS           64
Linux           62
Windows 7       45
Chrome OS       27
macOS           12
Mac OS X        8
Windows 10 S    8
Android         2
Name: count, dtype: int64
```

```
[233] : def cat_os(input):      # cat_os means category of OS
        if input == 'Windows 10' or input == 'Windows 7' or input == 'Windows 10 S':
            return 'Windows'
        elif input == 'macOS' or input == 'Mac OS X':
            return 'Mac'
        else:
            return 'Other OS/Linux/No OS'
```

```
[234] : df['os'] = df['OpSys'].apply(cat_os)
```

```
[235] : df.head()
```

```
[235]: Company  TypeName  Ram  OpSys  Weight      Price  Touchscreen  IPS  \
0   Apple  Ultrabook   8  macOS   1.37  71378.6832         0      1
1   Apple  Ultrabook   8  macOS   1.34  47895.5232         0      0
2    HP    Notebook   8  No OS   1.86  30636.0000         0      0
3   Apple  Ultrabook  16  macOS   1.83  135195.3360         0      1
4   Apple  Ultrabook   8  macOS   1.37  96095.8080         0      1
```

```
      ppi      Cpu brand  Gpu brand      os
0  226.983005 Intel  Core i5      Intel      Mac
1  127.677940 Intel  Core i5      Intel      Mac
2  141.211998 Intel  Core i5      Intel  Other OS/Linux/No OS
3  220.534624 Intel  Core i7      AMD      Mac
4  226.983005 Intel  Core i5      Intel      Mac
```

```
[236] : df.drop(columns=['OpSys'], inplace=True)
```

```
[237] : df.head()
```

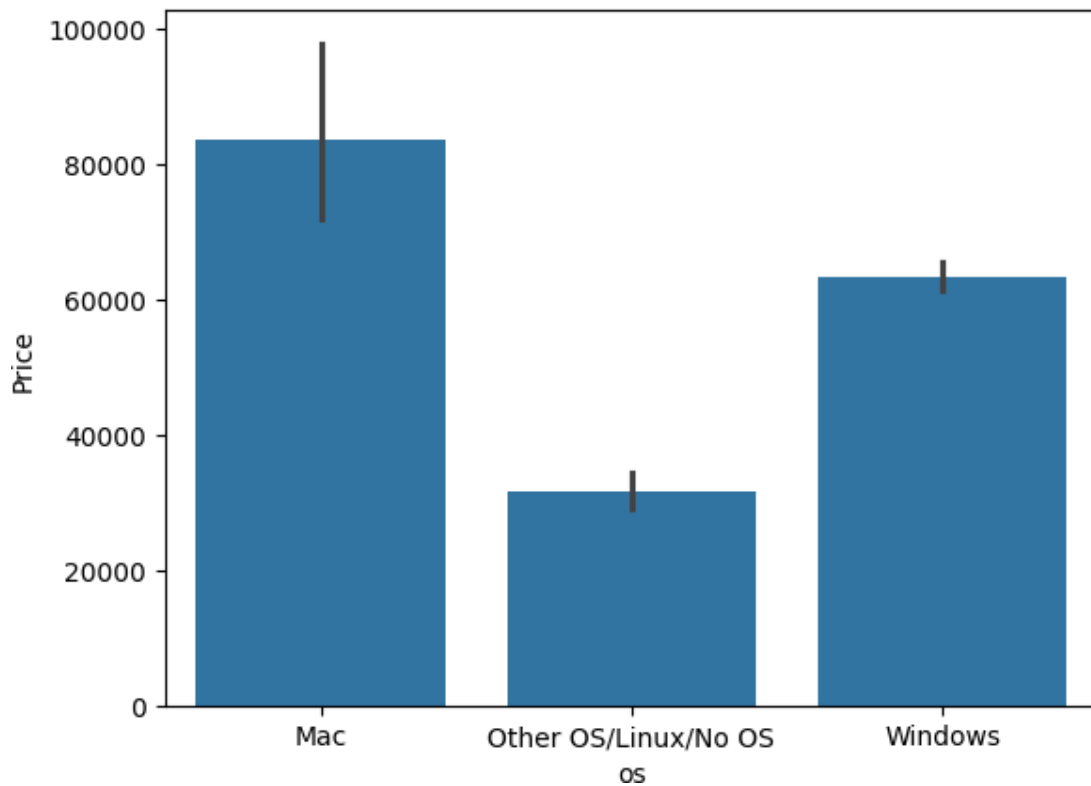
```
[237] : Company  TypeName  Ram  Weight      Price  Touchscreen  IPS      ppi  \
0   Apple  Ultrabook   8    1.37  71378.6832         0      1  226.983005
1   Apple  Ultrabook   8    1.34  47895.5232         0      0  127.677940
```

2	HP	Notebook	8	1.86	30636.0000	0	0	141.211998
3	Apple	Ultrabook	16	1.83	135195.3360	0	1	220.534624
4	Apple	Ultrabook	8	1.37	96095.8080	0	1	226.983005

	Cpu brand	Gpu brand	os
0	Intel Core i5	Intel	Mac
1	Intel Core i5	Intel	Mac
2	Intel Core i5	Intel	Other OS/Linux/No OS
3	Intel Core i7	AMD	Mac
4	Intel Core i5	Intel	Mac

```
[238] : sns.barplot(x=df['os'], y=df['Price'])
```

```
[238] : <Axes: xlabel='os', ylabel='Price'>
```



Weight

```
[239] : sns.distplot(df['Weight'])
```

<ipython-input-239-05ee4c8848be>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Weight'])
```

[239] : <Axes: xlabel='Weight', ylabel='Density'>



[240] : df.dtypes

```
[240] : Company      object
      TypeName      object
      Ram           int32
      Weight        float32
      Price         float64
      Touchscreen    int64
      IPS           int64
      ppi           float64
```

```
Cpu brand      object
Gpu brand      object
os             object
dtype: object
```

```
[241] : sns.distplot(df['Price'])      # We can see the graph is skewed
```

```
<ipython-input-241-0a9ec3ba582e>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Price'])      # We can see the graph is skewed
```

```
[241] : <Axes: xlabel='Price', ylabel='Density'>
```



```
[242] : sns.distplot(np.log(df['Price']))      # Normalized graph
```

<ipython-input-242-cfba2a5ffa37>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(np.log(df['Price']))      # Normalized graph
```

```
[242] : <Axes: xlabel='Price', ylabel='Density'>
```



```
[243] : X = df.drop(columns=['Price'])
        y = np.log(df['Price'])      # To normalise the data
```

```
[244] : X
```



```
[244] :
```

	Company		TypeName	Ram	Weight	Touchscreen	IPS	ppi	\
0	Apple		Ultrabook	8	1.37	0	1	226.983005	
1	Apple		Ultrabook	8	1.34	0	0	127.677940	
2	HP		Notebook	8	1.86	0	0	141.211998	
3	Apple		Ultrabook	16	1.83	0	1	220.534624	
4	Apple		Ultrabook	8	1.37	0	1	226.983005	
...	
1301	Lenovo	2 in 1	Convertible	4	1.80	1	1	157.350512	
1302	Lenovo	2 in 1	Convertible	16	1.30	1	1	276.053530	
1303	Lenovo		Notebook	2	1.50	0	0	111.935204	
1304	HP		Notebook	6	2.19	0	0	100.454670	
1305	Asus		Notebook	4	2.20	0	0	100.454670	

	Cpu brand	Gpu brand	os
0	Intel Core i5	Intel	Mac
1	Intel Core i5	Intel	Mac
2	Intel Core i5	Intel	Other OS/Linux/No OS
3	Intel Core i7	AMD	Mac
4	Intel Core i5	Intel	Mac
...
1301	Intel Core i7	Intel	Windows
1302	Intel Core i7	Intel	Windows
1303	Other Intel Processor	Intel	Windows
1304	Intel Core i7	AMD	Windows
1305	Other Intel Processor	Intel	Windows

[1296 rows x 10 columns]

```
[245] :
```

```
y
```

```
[245]:
```

0	11.175755
1	10.776777
2	10.329931
3	11.814476
4	11.473101

```
...
```

1301	10.433899
1302	11.288115
1303	9.409283
1304	10.614129
1305	9.886358

Name: Price, Length: 1296, dtype: float64

```
[246] :
```


```

from sklearn.compose import ColumnTransformer
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import OneHotEncoder

```

```
[260]: sns.heatmap(df_encoded.corr())
```

```
[260]: <Axes: >
```



```
[247] : df_encoded = pd.get_dummies(df, columns=['Company', 'TypeName', 'Cpu brand', '_
        'Gpu brand', 'os'])
```

```
[248] : df_encoded
```

```
[248]:
```

	Ram	Weight	Price	Touchscreen	IPS	ppi	Company_Acer \
0	8	1.37	71378.6832	0	1	226.983005	False
1	8	1.34	47895.5232	0	0	127.677940	False
2	8	1.86	30636.0000	0	0	141.211998	False
3	16	1.83	135195.3360	0	1	220.534624	False
4	8	1.37	96095.8080	0	1	226.983005	False
...
1301	4	1.80	33992.6400	1	1	157.350512	False
1302	16	1.30	79866.7200	1	1	276.053530	False

1303	2	1.50	12201.1200	0	0	111.935204	False
1304	6	2.19	40705.9200	0	0	100.454670	False
1305	4	2.20	19660.3200	0	0	100.454670	False

	Company_Apple	Company_Asus	Company_Chuiwi	...	\
0	True	False	False	...	
1	True	False	False	...	
2	False	False	False	...	
3	True	False	False	...	
4	True	False	False	...	
...	
1301	False	False	False	...	
1302	False	False	False	...	
1303	False	False	False	...	
1304	False	False	False	...	
1305	False	True	False	...	

	Cpu brand_Intel Core i5	Cpu brand_Intel Core i7	\
0	True	False	
1	True	False	
2	True	False	
3	False	True	
4	True	False	
...	
1301	False	True	
1302	False	True	
1303	False	False	
1304	False	True	
1305	False	False	

	Cpu brand_Other Intel Processor	Gpu brand_AMD	Gpu brand_ARM	\
0	False	False	False	
1	False	False	False	
2	False	False	False	
3	False	True	False	
4	False	False	False	
...	
1301	False	False	False	
1302	False	False	False	
1303	True	False	False	
1304	False	True	False	
1305	True	False	False	

	Gpu brand_Intel	Gpu brand_Nvidia	os_Mac	os_Other	OS/Linux/No OS	\
0	True	False	True		False	
1	True	False	True		False	
2	True	False	False		True	

3	False	False	True	False
4	True	False	True	False
...
1301	True	False	False	False
1302	True	False	False	False
1303	True	False	False	False
1304	False	False	False	False
1305	True	False	False	False

	os_Windows
0	False
1	False
2	False
3	False
4	False
...	...
1301	True
1302	True
1303	True
1304	True
1305	True

[1296 rows x 43 columns]

```
[249]: df_encoded.columns
```

```
[249] : Index(['Ram', 'Weight', 'Price', 'Touchscreen', 'IPS', 'ppi', 'Company_Acer',
      'Company_Apple', 'Company_Asus', 'Company_Chuiwi', 'Company_Dell',
      'Company_Fujitsu', 'Company_Google', 'Company_HP', 'Company_Huawei',
      'Company_LG', 'Company_Lenovo', 'Company_MSI', 'Company_Mediacom',
      'Company_Microsoft', 'Company_Razer', 'Company_Samsung',
      'Company_Toshiba', 'Company_Vero', 'Company_Xiaomi',
      'TypeName_2 in 1 Convertible', 'TypeName_Gaming', 'TypeName_Netbook',
      'TypeName_Notebook', 'TypeName_Ultrabook', 'TypeName_Workstation',
      'Cpu_brand_AMD Processor', 'Cpu_brand_Intel Core i3',
      'Cpu_brand_Intel Core i5', 'Cpu_brand_Intel Core i7',
      'Cpu_brand_Other Intel Processor', 'Gpu_brand_AMD', 'Gpu_brand_ARM',
      'Gpu_brand_Intel', 'Gpu_brand_Nvidia', 'os_Mac',
      'os_Other OS/Linux/No OS', 'os_Windows'],
      dtype='object')
```

```
[250] : X = df_encoded.drop('Price', axis=1) # 'Price' is your target variable
      y = np.log(df_encoded['Price'])
```

```
[251] : X
```

[251]:

	Ram	Weight	Touchscreen	IPS	ppi	Company_Acer	Company_Apple	\
0	8	1.37	0	1	226.983005	False	True	
1	8	1.34	0	0	127.677940	False	True	
2	8	1.86	0	0	141.211998	False	False	
3	16	1.83	0	1	220.534624	False	True	
4	8	1.37	0	1	226.983005	False	True	
...	
1301	4	1.80	1	1	157.350512	False	False	
1302	16	1.30	1	1	276.053530	False	False	
1303	2	1.50	0	0	111.935204	False	False	
1304	6	2.19	0	0	100.454670	False	False	
1305	4	2.20	0	0	100.454670	False	False	

Company_Asus	Company_Chuiwi	Company_Dell	...	Cpu brand_Intel	Core i5	\
False	False	False	...		True	
False	False	False	...		True	
False	False	False	...		True	
False	False	False	...		False	
False	False	False	...		True	
...	
False	False	False	...		False	
False	False	False	...		False	
False	False	False	...		False	
False	False	False	...		False	
True	False	False	...		False	

Cpu brand_Intel	Core i7	Cpu brand_Other	Intel	Processor	Gpu brand_AMD	\
	False			False	False	
	False			False	False	
	False			False	False	
	True			False	True	
	False			False	False	
...	
	True			False	False	
	True			False	False	
	False			True	False	
	True			False	True	
	False			True	False	

Gpu brand_ARM	Gpu brand_Intel	Gpu brand_Nvidia	os_Mac	\
False	True	False	True	
False	True	False	True	
False	True	False	False	
False	False	False	True	
False	True	False	True	
...	
False	True	False	False	

1302	False	True	False	False
1303	False	True	False	False
1304	False	False	False	False
1305	False	True	False	False

	os_Other	OS/Linux/No OS	os_Windows
0		False	False
1		False	False
2		True	False
3		False	False
4		False	False
...	
1301		False	True
1302		False	True
1303		False	True
1304		False	True
1305		False	True

[1296 rows x 42 columns]

[252] :

y

[252]:

0	11.175755
1	10.776777
2	10.329931
3	11.814476
4	11.473101

...	
1301	10.433899
1302	11.288115
1303	9.409283
1304	10.614129
1305	9.886358

Name: Price, Length: 1296, dtype: float64

[253] : X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

[254] : rf_model = RandomForestRegressor(random_state=42)
rf_model.fit(X_train, y_train)

[254] : RandomForestRegressor(random_state=42)

[255] : y_pred = rf_model.predict(X_test)

[256] : # Evaluate the model
mse = mean_squared_error(y_test, y_pred)

```
print("Mean Squared Error:", mse)
```

Mean Squared Error: 0.05299513432325618

```
[257] : # Calculate RMSE  
rmse = np.sqrt(mse)  
print("Root Mean Squared Error (RMSE):", rmse)
```

Root Mean Squared Error (RMSE): 0.23020672084727714

```
[258] : from sklearn.metrics import r2_score  
  
# Calculate R^2 score  
r2 = r2_score(y_test, y_pred)  
  
print("R^2 Score:", r2)
```

R^2 Score: 0.8570143929282722